

Image-Related Neural Networks

How to see like a human

Yordan Darakchiev

Technical Trainer

iordan93@gmail.com





sli.do

#DeepLearning

Table of Contents

- Convolutional neural networks
 - Operations
 - Architectures
- Generalizations
 - ResNet
 - 1x1 convolutions, "network-in-network"
 - Object localization

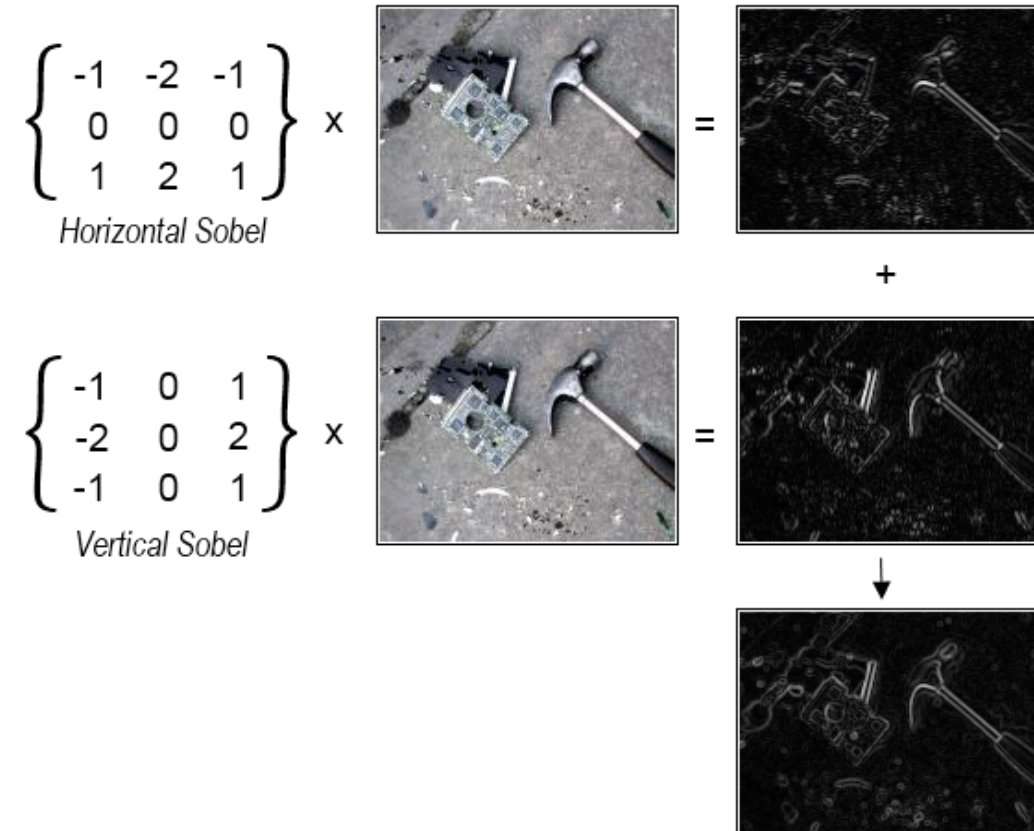


Convolutional Neural Networks

Learning from images

Convolution

- Given an image I and a filter F , $R = I \circledast F$ is defined as
 - For each pixel (i, j) , $R_{ij} = \sum(I_{ij} * F)$
- Depending on F , the result has different meanings
 - Example: Sobel edge detection
- F is usually square, with odd rank (so that it has a central pixel)
 - E.g. 3, 5, 7



Convolution (2)

- Padding

- "Valid convolution": no padding
- "Same convolution": pad so that the output size remains unchanged:
$$p = \frac{f-1}{2}, f - \text{filter size}$$

- Sliding window: **stride** s

- How many pixels we should skip


- Summary

- Input


- $n \times n$ image
- $f \times f$ filter
- padding p
- stride s

- Output image dimensions: $\left\lfloor \frac{n+2p-f}{s} + 1 \right\rfloor \times \left\lfloor \frac{n+2p-f}{s} + 1 \right\rfloor$
- If the image is non-square, adjust the dimensions in the formula


7 x 7 Input Volume




5 x 5 Output Volume



7 x 7 Input Volume

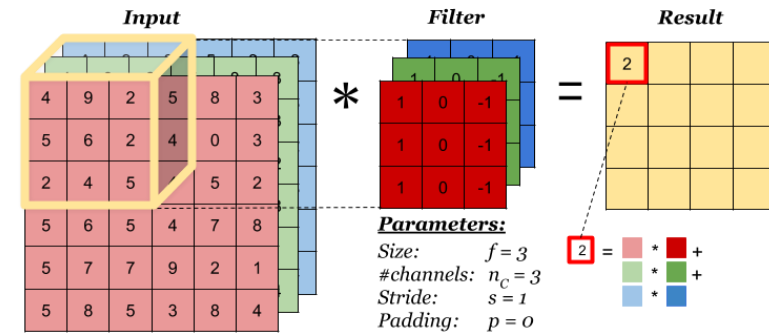


3 x 3 Output Volume




Convolution over Volume

- If the image has many channels with dimensions $n \times n \times c$, just use an $f \times f \times c$ filter
 - I.e. apply the operation independently for each channel
 - Result: 2D image
- Many filters
 - Each one produces a 2D image
 - Stack them together (since they're independent)
⇒ 3D volume
- The convolution operations we use operate over 3D volumes




Convolutional Layers

- Just like a regular network
 - Input volume dimensions $n \times n \times c$
 - Choose f , n_f (number of filters), p , s
 - Learn each value of the filters, apply bias terms
 - Add non-linearity (e.g. ReLU)
 - Convolutions are linear operations
 - Sometimes, the convolution and activation layers are shown separately
 - Produce output volume



Pooling

- Used to reduce the number of parameters in the next layers
- Applied like convolution
- Parameters: window size f , stride s , operation
 - Most commonly used operation: max ([max-pooling](#))
 - In the past: [avg-pooling](#) was also widely used
 - Other operations are possible but uncommon
- [No trainable parameters](#)



Why Do Convolutions Work?

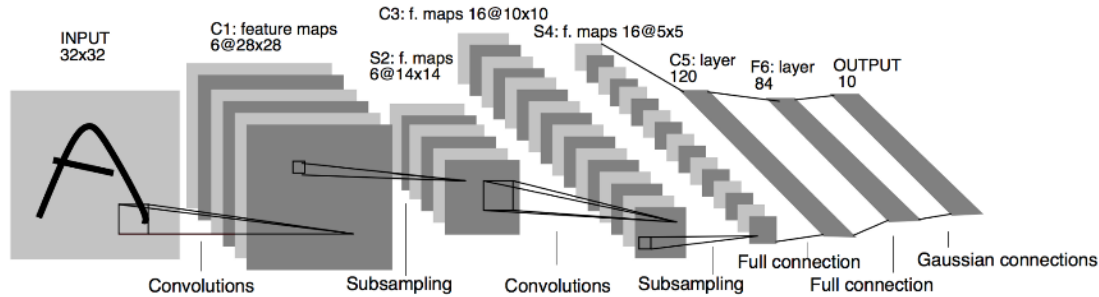
- Image assumptions
 - Individual features are relatively localized
 - The relative (not absolute) position of features is really important
- Convolutions help us to **share computations**
 - An edge detector is useful in many parts of the image
- Each filter has a **low-dimensional input**
 - Simplifies computations
- Visualizing and Understanding Convolutional Networks, Matthew Zeiler, 2014

Convolutional Layer Architecture

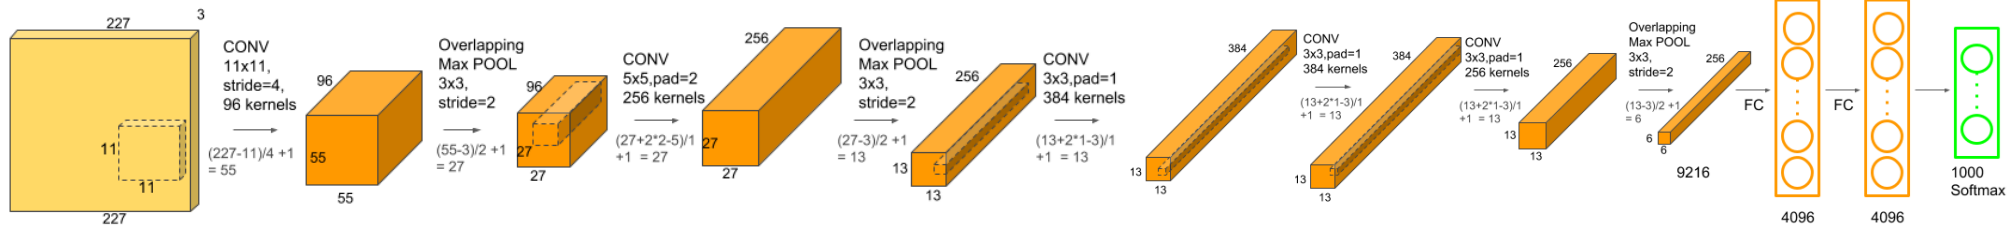
- Input volume: $h \times w \times c$
- Parameters: f, p, s, n_f
- $h' = \left\lfloor \frac{h+2p-f}{s} + 1 \right\rfloor, w' = \left\lfloor \frac{w+2p-f}{s} + 1 \right\rfloor$
- Filter dimensions: $f \times f \times c$, total of n_f filters
- Weights, biases: like fully-connected layers
 - $W = f \times f \times c \times n_f, b = 1 \times 1 \times 1 \times n_f$
- After computing, apply activation function
- Output volume: $h' \times w' \times n_f$

Convolutional Neural Networks

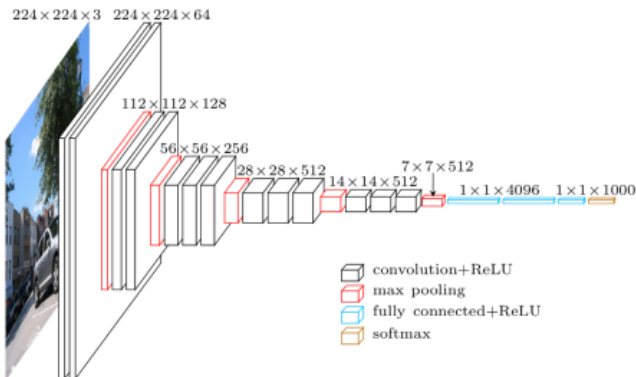
■ LeNet-5 ([Yann LeCun, 1998](#))



■ AlexNet ([Alex Krizhevsky, 2012](#))



■ VGG-19 ([Karen Simonyan, 2014](#))








Generalizations and Expansions

Applying other tricks

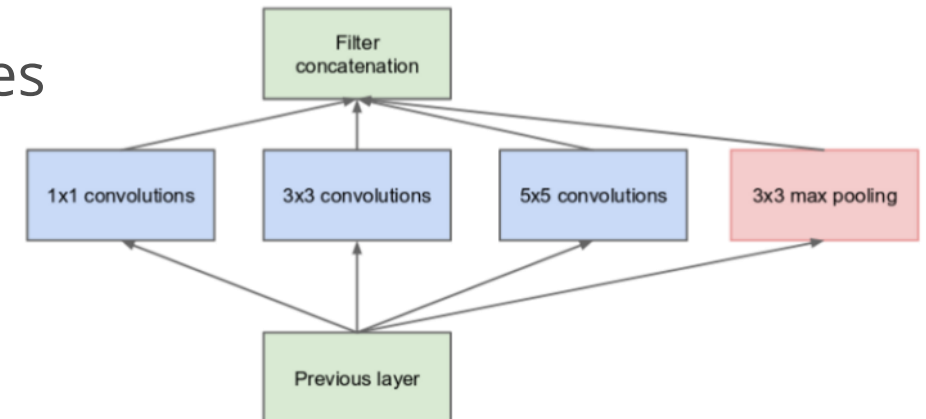
Residual Networks (ResNets)

- Deeper networks allow us to compute complex functions
 - But are difficult to train (e.g. vanishing gradients)
 - The validation error increases (not by overfitting)
- Solution: shortcut connections
 - Pass the activation skipping 1 or more layers
 - Reason: the identity function $y = x$ is really easy to learn
- Results
 - ImageNet
 - 18 / 34 layers
 - Runtime: faster than VGG



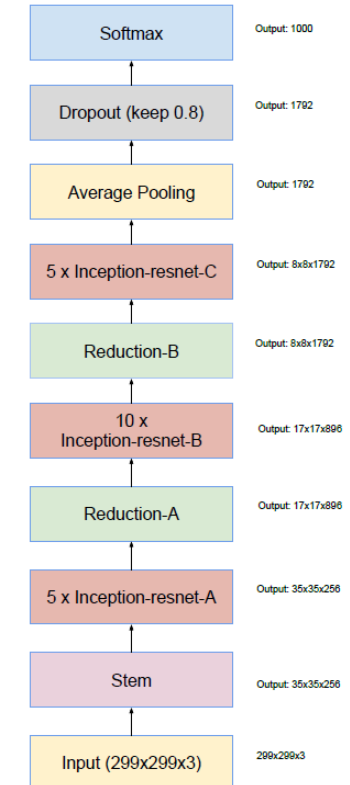
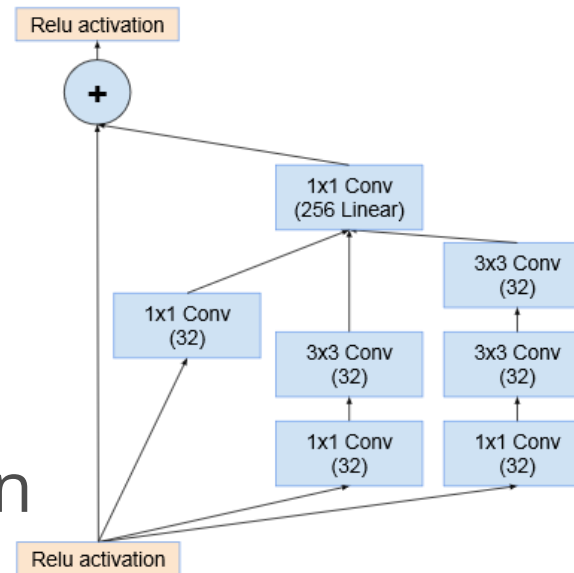
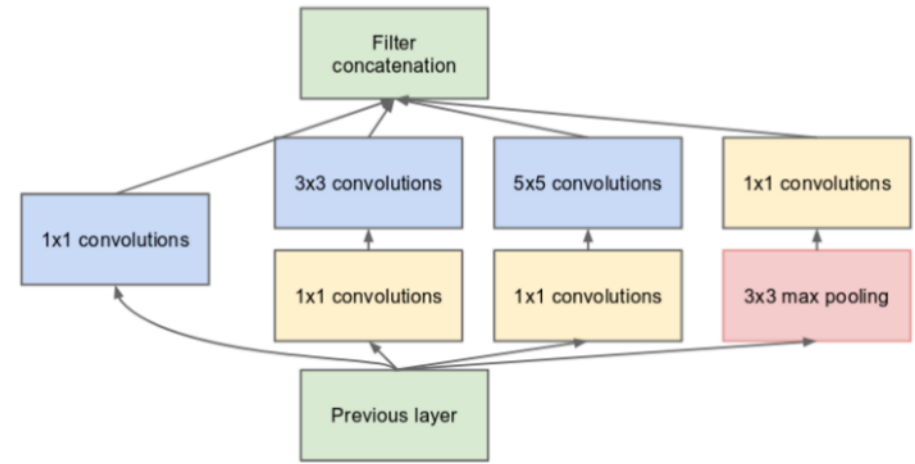
1x1 Convolutions

- Network in Network: Min Lin, 2013
- Single filter: does nothing (scales the input image)
- Many filters: $(w \times h \times c) \circledast (1 \times 1 \times c \times n_f) = (w \times h \times n_f)$
 - Keeps the image dimensions, changes the third dimension
 - Example: $(28 \times 28 \times 192)$, 32 filters $(1 \times 1 \times 192) \Rightarrow (28 \times 28 \times 32)$
 - Dimensionality reduction
- Inception (v1, v2 and v3, v4 and Inception-Res-Net)
 - Main idea
 - Kernel size corresponds to "size" of features
 - We can't know the kernel size f , so try many sizes and let the network decide what's most useful





Inception

- Problem: lots of computations
 - Solution: dimensionality reduction before each convolution
 - "Inception block"
- GoogLeNet: 9 inception modules
- Inception-Res-Net
 - A simple combination of the two concepts
 - Idea: create a deeper inception block, simplifying learning through a ResNet connection



Object Localization

- Input: image; output: bounding box (x, y, w, h)
 - Regression
 - Classification and localization
 - Simplest case: 1 object
 - Output a vector: $[p, x, y, w, h, c_1, c_2, \dots, c_k]$
 - $p = 0 \Rightarrow$ no object detected; we don't care about the other numbers
 - $p = 1 \Rightarrow$ object detected; class: c_1, \dots, c_k ; bounding box x, y, w, h
 - Metrics: usually IoU (or Euclidean distance)
 - Implementations: YOLO (You Only Look Once)
 - Also: R-CNN (Region-proposing network)
- 
- A photograph of a city street scene with various objects labeled with bounding boxes and class names. The labels include 'car', 'bus', 'traffic light', 'person', and 'handbag'. The bounding boxes are colored in yellow, pink, blue, and green. The scene shows a multi-lane road with cars, a bus, and pedestrians, with buildings and a traffic light in the background.



Summary

- Convolutional neural networks
 - Operations
 - Architectures
- Generalizations
 - ResNet
 - 1x1 convolutions, "network-in-network"
 - Object localization

The image features a white background with two thick, wavy blue bars at the top and bottom. The top bar is a lighter blue, while the bottom bar is a darker blue. Centered on the white background is the word "Questions?" in a large, bold, blue sans-serif font.

Questions?