

# Natural Language Processing

Taking time into account

**Yordan Darakchiev**

Technical Trainer

[iordan93@gmail.com](mailto:iordan93@gmail.com)





sli.do

#DeepLearning

# Table of Contents

- Time-dependent (sequential) models
  - Architecture
  - Types
- Improvements
- Word (token) representations
- Refinement algorithms

# Time-Dependent Models

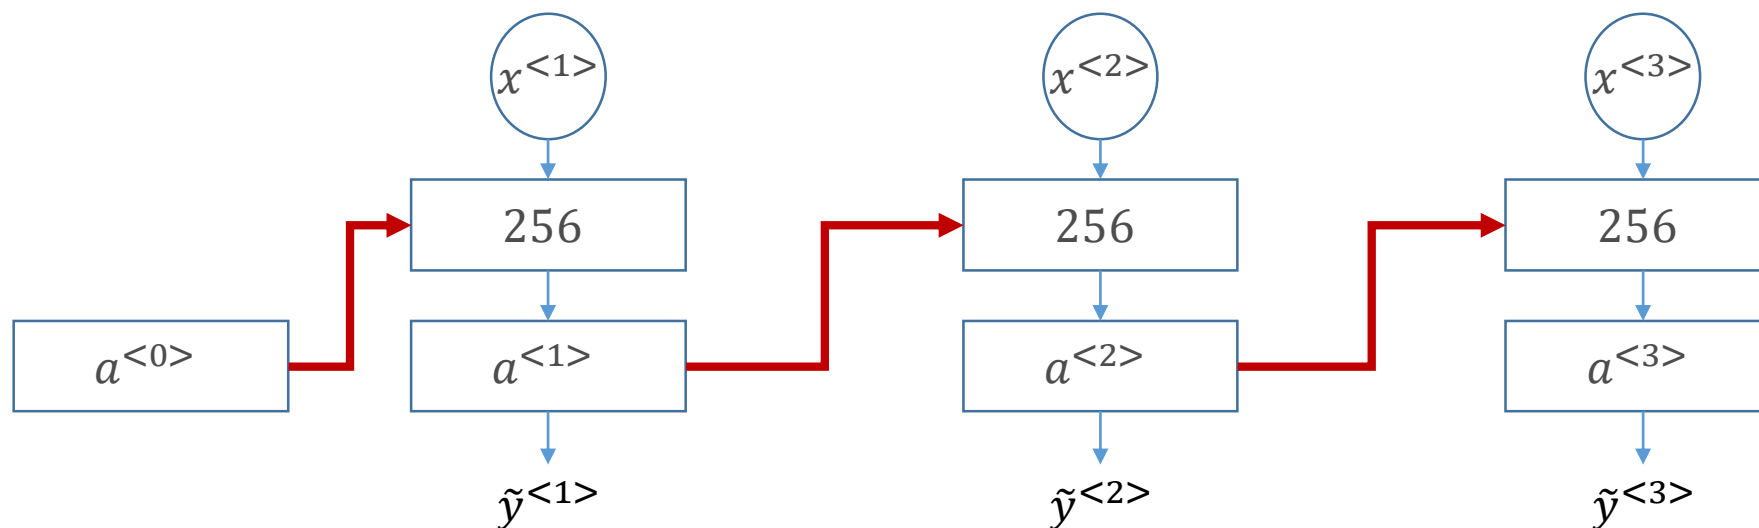
**"Time is precious, waste it wisely"**

# Time-Dependent Model Examples

- Speech recognition: audio → transcript
- Machine translation: text (EN) → text (FR)
- Activity recognition: video → activity type (e.g. walking)
- Sentiment analysis: text → sentiment
- Generation
  - Text summarization
  - Music generation
- More generally, models whose inputs depend on time
  - "Standard" models:  $\tilde{y} = f(x)$ ; recurrent models:  $\tilde{y} = f(x, s)$ 
    - $s$  – current state
  - Standard models don't allow variable-length inputs
  - Most standard models don't allow for weight sharing


# Working with Sequences

- Training example:  $x = \text{"A black cat in a box"}$ 
  - Split words (**tokenize** the input)
  - Present words as 1-hot encoded vectors using a dictionary (**vocabulary**)
  - $x^{<1>} = \text{"a"} = [1 \ 0 \ \dots \ 0]^T \equiv V_1$
  - $x^{<2>} = \text{"black"} = [0 \ 0 \ \dots \ 1 \ \dots \ 0]^T \equiv V_{329}$ , etc.
- Take a standard model (1-layer NN), pass each word



# Recurrent Neural Networks

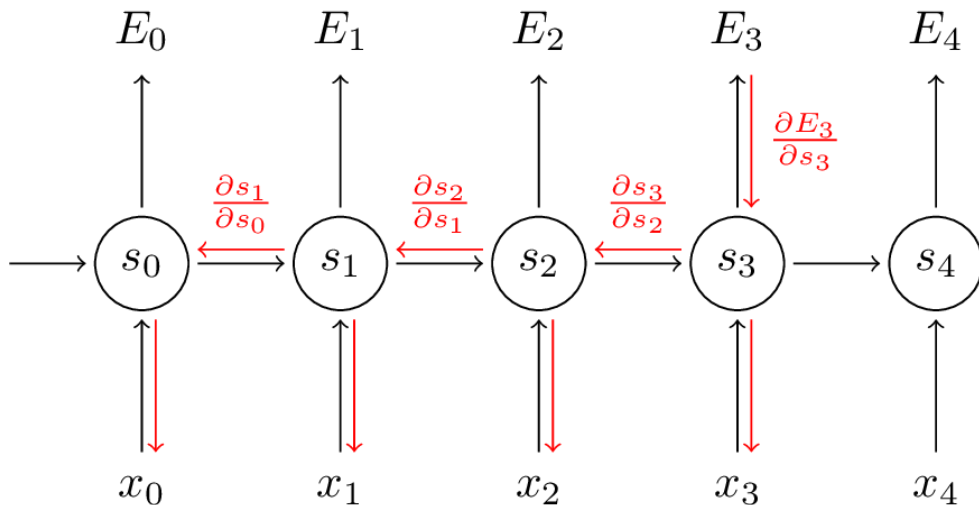

- "RNN cell"



- Deep architectures are also possible


- Learning: backpropagation through time

- The same as in a multi-layer network:  $\min J(y^{<t>}, \tilde{y}^{<t>})$



# RNN Architectures

- One to one: standard
- One to many
  - Sequence generation given seed (e.g. image captioning)
- Many to one
  - One output for sequence (e.g. sentiment analysis)
- Many to many
  - Encoding and decoding (e.g. machine translation)
  - Synchronized output (e.g. video classification for each frame)





# Language Model

## ■ Training

- Tokenize the input  $x = [x^{<1>}; x^{<2>}; x^{<T_x>}]$
- Use a standard RNN, with no initial seed
  - $a^{<0>} = [0 \ 0 \ \dots \ 0] = \vec{0}, x^{<0>} = \vec{0}$
  - Output:  $\tilde{y}$ : a vector of probabilities for each word [0,0385 0,0476 ... 0,00041]
    - Softmax, with 10 000 outputs

## ■ Explanation

- First token:  $\tilde{y} = P(w_1)$
- Second token:  $\tilde{y} = P(w_2|w_1)$
- In general:  $\tilde{y} = P(w_k|w_1, w_2, \dots, w_{k-1})$

## ■ Generation: random sampling according to computed $P$


- Input  $x^{<0>} = \vec{0}, a^{<0>} = \vec{0}$ ; compute  $a^{<1>}, \tilde{y}^{<1>}$ ; choose a word  $w_1$
- Input  $x^{<1>} \equiv w_1, a^{<1>}$ ; compute  $a^{<2>}, \tilde{y}^{<2>}$ ; choose a word  $w_2$
- ... until you reach [.]

# Improved Models

**Making things even more difficult**

# Vanishing Gradients

- RNN with a long input is similar to a very deep NN
- Examples
  - The match was long, but we won it which made us happy.
  - We **decided to go to the movies**, but our friend, who doesn't like scary movies, **didn't want to go**.
- Solution: Gated recurrent unit (GRU) – [Cho et al., 2014](#)
  - Update gate ( $z_t$ ): how much of the past information to retain
  - Reset gate ( $r_t$ ): how much information to forget
  - Final memory: current information + previous "context"



$$z_t = \sigma(W_z \cdot [h_{t-1}, x_t])$$


$$r_t = \sigma(W_r \cdot [h_{t-1}, x_t])$$

$$\tilde{h}_t = \tanh(W \cdot [r_t * h_{t-1}, x_t])$$

$$h_t = (1 - z_t) * h_{t-1} + z_t * \tilde{h}_t$$

# Long-Short Term Memory (LSTM)


- Even more powerful (and complicated)
  - [Hochreiter and Schmidhuber, 1997](#)
  - This is only one layer, LSTM layers can also be stacked



- Basic parts of the architecture
  - Forget gate  $f_t$
  - Update gate  $i_t$
  - Cell state  $C_t$
  - Output  $o_t$
- A good [article](#) explaining LSTM cells

# Bi-Directional Networks

- Intuition
  - RNNs may need information ahead, "from the future"
  - E.g. to translate word  $w^{<5>}$  we may need the whole sentence
- Solution: just create pairs of networks



- These can be RNN, GRU, LSTM or other layers
- To compute activations, go left to right, then right to left

# Representing Words

**Find your way  
in the multi-dimensional space**

# Word Representation

- Basic idea: one-hot encoding
- How to get insights on word relations?
  - Try to estimate **word features**: vectors of numbers for each word
    - Unsupervised process
    - **Embedding** from a space with one dimension per word to a lower-dimensional space (e.g. 300D)
  - Example uses
    - Use similarity measures (e.g. cosine distance) between vectors
    - Use projections to generate analogies ([Mikolov et al., 2013](#))
- Visualization: usually **t-SNE** or PCA
- Tensorboard uses Google's Embedding Projector
  - <https://www.tensorflow.org/guide/embedding>

# Word2Vec and GloVe

- What we already described
  - A matrix  $E$  where each vocabulary word has a dense vector
- Context-target word pairs
  - Compute vectors for context and target
  - Loss: cross-entropy
- Similarity
  - Cosine similarity; closest words to "Sweden"
- Associations
  - Rome : Italy :: Beijing : China
  - king : queen :: man : woman
  - [Other examples](#)

Word	Cosine distance
norway	0.760124
denmark	0.715460
finland	0.620022
switzerland	0.588132
belgium	0.585835
netherlands	0.574631
iceland	0.562368
estonia	0.547621
slovenia	0.531408



# Refinement Algorithms

Some more tricks up our sleeves

# Beam Search

- Translation
  - Similar to generation,  $\tilde{y} = f(x)$ , maximize  $P(\tilde{y}|x)$
- What if we have multiple candidates?
  - Use the language model to compute  $P$
- Slight complications
  - *I am visiting NY this year end.*
  - *I am going to be visiting NY this year end.*
  - $P(\text{going}|\text{i, am}) > P(\text{visiting}|\text{i, am})$
  - Observations
    - One word at a time doesn't work too well
    - All words will require enormous computation power
  - Solution: **Beam search**
    - At each step, choose top  $B$  words (**beam width**)
    - More [details](#)

# Attention

- [Xu et al., 2015](#)
- Another mechanism for dealing with complicated inputs
  - Another caveat: longer sentences have inherently lower probabilities so models tend to favor short sentences
  - Intuition: we don't need to know the entire sequence in order to be able to translate
- Idea
  - Use a bi-directional RNN (or GRU / LSTM)
  - For each part of the input  $x^{<t>}$ , compute "how much you care" about it: *attention*<sup><t></sup>
- [Usages](#)
  - Translation, image captioning, speech recognition, text summarization, etc.

# Summary

- Time-dependent (sequential) models
  - Architecture
  - Types
- Improvements
- Word (token) representations
- Refinement algorithms

The image features a white background with two decorative blue bars. The top bar is a solid blue strip. Below it is a white space containing the text. At the bottom, there is another white space, followed by a thin dark blue line and a final solid blue strip at the very bottom.

Questions?